

# TLS For IoT

Illya Gerasymchuk  
Instituto Superior Tecnico  
Lisbon, Portugal  
illya@iluxonchik.me

**Abstract**—Transport Layer Security (TLS) is, one of the most used communication security protocols, it is however, not suitable in the context of Internet Of Things (IoT). The resource-limited nature of a big part of IoT devices does not allow for the use of computationally complex and memory demanding operations present in a standard TLS implementation. Most of previous work focused entirely on Datagram TLS (DTLS) and can not be easily integrated with existing deployments. This work focuses on how TLS and the extension mechanism can be used to define a framework to adapt the protocol to specific needs. This is the approach that will be followed in the second part of the work. Having an adaptable and easy to use solution is crucial for its adaptation in IoT, where security might have been completely foregone otherwise.

**Index Terms**—TLS, DTLS, SSL, IoT, cryptography, protocol, lightweight cryptography

## I. INTRODUCTION

The Internet of Things (IoT) is a network of devices, from simple sensors to smartphones and wearables which are connected together. In fact, it can be any other object that has an assigned IP address and is provided with the ability to transfer data over a network. Even a salt shaker [1] can now be part of the global network.

The IoT technology provides many benefits, from personal comfort to transforming entire industries, mainly due to increased connectivity and new sources for data analysis. The technological development, however, tends to focus on innovative design rather than on privacy and security. IoT devices frequently connect to networks using inadequate security and are hard to update when vulnerabilities are found.

This lack of security in the IoT ecosystem has been exploited by the the *Mirai* botnet [2] when it overwhelmed several high-profile targets with massive Distributed Denial-Of-Service (DDoS) attacks. This is the most devastating attack involving IoT devices done to date. However, the *Reaper* botnet [3] could be even worse if it is ever put to malicious use. Similar attacks will inadvertently come in the future.

TLS is one of the most used security protocols in the world, allowing two peers to communicate securely. It is designed to run on top of a reliable, connection-oriented protocol, such as TCP. Datagram TLS (DTLS) is the version of TLS that runs on top of an unreliable transport protocol, such as UDP. Most IoT devices have very limited processing power, storage and energy. Moreover, the performance of TCP is known to be inefficient in wireless networks, due to its congestion control algorithm. This situation is worsened with the use of low-power radios and lossy links found in sensor networks.

Therefore, the use of TCP with IoT is usually not the best option. For this reason, DTLS, which runs on top of UDP, is used more frequently in such devices. The work that will be done in the context of this dissertation, can however, be applied to either one of them, so even though mostly TLS will be mentioned, almost everything can also be applied to DTLS. This is a consequence of DTLS being just an adaption of TLS over unreliable transport protocols, with no changes done to the core protocol.

The problem in using (D)TLS in IoT is that it is not lightweight, since it has not been designed for such environments. An IoT device may only have 256 KB of RAM and needs to conserve the battery, while sending and receiving a large amount of small information constantly. For example, consider the case of a temperature sensor that sends temperature measures every 30 seconds to a server. In this case it just needs to send a few bytes of data and do it with minimal overhead, to conserve RAM and battery. If that sensor is going to use (D)TLS 1.2, it will need two extra roundtrips before it can send any data. This can result in an overhead of several hundreds of milliseconds. Besides that, it will need to perform heavy mathematical operations involved in cryptography, using even more energy and taking even more time. Given this, there is a clear need for a more lightweight (D)TLS for the IoT. This is the problem that we will be addressing in this work.

In the process of the work on this dissertation, we have made several contributions to the TLS 1.3 specification [4]. Among several contributions, we corrected some inconsistencies/incompletenesses in the specification (e.g. the document mentioned some extensions as being mandatory, yet they were not present in the "Mandatory-to-Implement" section [5]) and some parts of the specification that were wrong, like the "Backwards Compatibility" section [6]. As a result, we were formally recognized as contributors and our name will be in the final document specifying the TLS 1.3 protocol. We are in no way affiliated with the Internet Engineering Task Force (IETF) or any of the entities responsible for the development of the TLS 1.3 specification.

In our work, we are using the *mbedTLS* library [7], which provides a lightweight TLS implementation for the IoT devices. In the process of profiling the library and studying its code, we have identified two bugs. The first one is related to the use of *SHA-1*-signed certificates [8] in example programs that come with *mbedTLS*, which are not supported in the default configuration. We have provided a fix for that issue [8]. The second bug is both: a deviation from the TLS 1.2

protocol specification and a security issue. The problem is that ciphersuites that should only allow *RSA*-signed certificates, besides those, also allow *ECDSA*-signed ones. We have reported the issue and provided proof of concept code [9].

The document is organized as follows: Section II describes the TLS protocol. Section III summarizes the related work. Section IV describes the goal of this work. In Section V, the work current work is outlined. Finally, the conclusion is done in Section VI.

## II. THE TLS PROTOCOL

TLS is a **client-server** protocol that runs on top a **connection-oriented and reliable transport protocol**, such as **TCP**. Its main goal is to provide **privacy** and **integrity** between the two communicating peers. Privacy implies that a third party will not be able to read the data, while integrity means that a third party will not be able to alter the data.

In the TCP/IP Protocol Stack, TLS is placed between the **Transport** and **Application** layers. It is designed to simplify the establishment and use of secure communications from the application developer's standpoint. The developer's task is reduced to creating a "secure" connection (*i.e.* socket), instead of a "normal" one.

A secure communication established using TLS has two phases.

- 1) The communicating peers authenticate one to another and negotiate security parameters, such as the secret keys and the encryption algorithm. This occurs under the Handshake Protocol.
- 2) The communicating peers exchange cryptographically protected data using the security parameters negotiated during the Handshake Protocol. This occurs under the Record Protocol.

In order to achieve its goals, during the Handshake Protocol the client and the server exchange various messages. The message flow is depicted in Figure 1. \* indicates situation-dependent messages that are not always sent.

TLS provides the following **security services**:

- **authentication** - both, **peer entity** and **data origin** (or **integrity**) authentication.
  - peer entity authentication** - a peer has a guarantee that it is talking to certain entity, for example, *www.google.com*. This is achieved through the use of Asymmetrical Cryptography (AC), also known as Public Key Cryptography (PKC), (*e.g.* RSA and DSA) or **symmetric key cryptography**, using a Pre-Shared Key (PSK).
- **confidentiality** - the data transmitted between the communicating entities (the client and the server) is encrypted. Symmetric cryptography is used for data encryption (*e.g.*, AES).
- **integrity** (also called **data origin authentication**) - a peer can be sure that the data was not modified or forged, *i.e.*, there is a guarantee that the received data is coming from the expected entity. For example, a peer can be

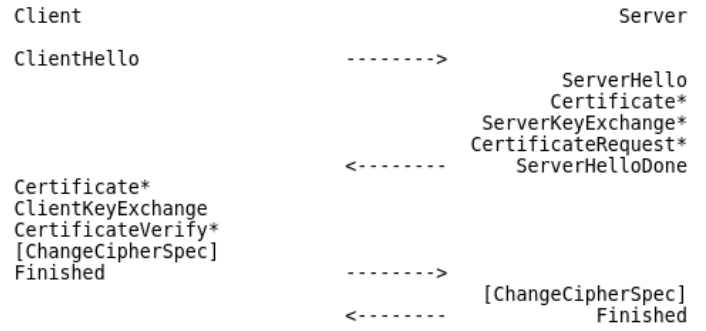


Fig. 1. TLS 1.2 message flow for a full handshake

sure that the *index.html* file that was sent to when it connected to *www.google.com* did, in fact, come from *www.google.com* and it was not tampered with by an attacker (**data integrity**). This is achieved either through the use of a keyed Message Authentication Code (MAC) or an Authenticated Encryption With Associated Data (AEAD) cipher.

- **replay protection** (also known as **freshness**) - a peer can be sure that a message has not been replayed. This is achieved through the use of sequence numbers. Each TLS record has a different sequence number, which is incremented. If a non-AEAD cipher is used, the sequence number is a direct input of the MAC function. If an AEAD cipher is used, a nonce derived from the sequence number is used as input to that cipher.
- **perfect forward secrecy (PFS)** - the confidentiality of the past interactions is preserved even if the long-term secret is compromised.

Despite using PKC, TLS does **not** provide **non-repudiation services**: neither **non-repudiation with proof of origin**, which addresses the peer denying the sending of a message, nor **non-repudiation with proof of delivery**, which addresses the peer denying the receipt of a message. This is due to the fact that instead of using **digital signatures**, either a keyed MAC or an AEAD cipher is used, both of which require a secret to be **shared** between the peers.

It is not required to use all of the security services every situation. In this sense, TLS is like a framework that allows to select which security services should be used for a communication session. As an example, certificate validation might be skipped, which means that the **authentication** guarantee is not provided. There are some differences regarding this claim between TLS 1.2 [10] and TLS 1.3 [11]. For example, while in the first there is a null cipher (no authentication, no confidentiality, no integrity), in the latter this is not true.

## III. RELATED WORK

Most of the work has been centered around DTLS, even though the majority of solutions can be applied to TLS as well. Herein we want to further explore TLS optimization. There is clearly a need for that, specially with Constrained Application Protocol (CoAP) over TCP and TLS standard being currently

developed [12]. The mentioned standard does not explore any TLS optimizations, and since any IoT device using it in the future would benefit from them, this is an important area to explore. CoAP [13] is often referred to as the "HTTP protocol for constrained devices".

None of the related work explored (D)TLS 1.3, mainly because the protocol is still in draft stages, however, major design changes are not expected at this point.

The majority of the work done in the area proposes a solution that is either tied to a specific protocol, such as CoAP, or requires an introduction of a third-party entity, such as the trust anchor in the case of the S3K system [14] or even both. This has two main issues. First, a protocol-specific solution cannot be easily used in an environment where (D)TLS is not used with that protocol. Second, the requirement of a third-party introduces additional cost and complexity, which will be a big resistance factor in adopting the technology. This is specially true for developers working on personal projects or projects for small businesses, leaving the communications insecure in the worse case scenario.

#### IV. PROPOSED SOLUTION

There is currently no tool which can suggest a (D)TLS configuration, based on the provided environment's limitations (e.g. available memory, power, processing speed) and security requirements (e.g. PFS, data privacy). The solution will be fully TLS compatible, since it will not modify any part of the protocol.

The goal of this work is to develop a framework that would suggest a (D)TLS configuration based on the needs and limitations of the environment. This configuration would include key exchange algorithms, encryption algorithms and TLS options (e.g. extensions), among others.

The solution will be developed for (D)TLS version 1.2, while also bearing in mind the new 1.3 version.

To achieve this goal, a thorough evaluation of every part of (D)TLS is needed. The main parts that will be evaluated are power consumption and the number CPU cycles used.

We would like to associate a quantitative cost with every (D)TLS feature. For example, we would like to answer questions such as "How much does Perfect Forward Secrecy cost, in terms of CPU cycles?". This will help the application developers in making decisions related to the security/performance trade off.

#### V. CURRENT WORK

In order to understand which features of (D)TLS should be used, as an answer to the limitations and requirements of the context, it's important to associate a cost (e.g. number of CPU cycles used) with each one of those features. For this reason, we are currently working on profiling the various parts of the TLS protocol.

All of the evaluation work is being done on the TLS implementation of the *mbedtls* library [15].

*mbedtls* has various TLS configurations. Some of those configurations are ciphersuites (key exchange algorithm +

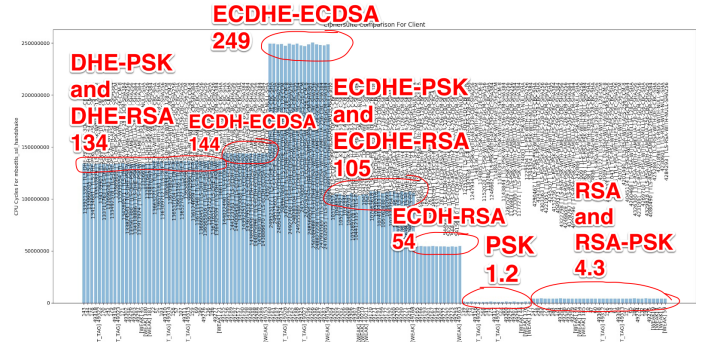


Fig. 2. Client ciphersuite profiling

encryption function + hash function), others are protocol features (e.g. session resumption, which allows to establish a connection using less computing resources, by using constructs from a previously established session), while others are extensions to the TLS protocol.

Currently, we are focused on profiling the TLS ciphersuites. More specifically, we are profiling the handshake portion of the protocol. After the handshake, the data is exchanged using the negotiated encryption and hash functions, as well as the associated secret keys. The latter is not TLS specific and comes down to evaluating the performance of various encryption and digest algorithms, which has already been done in various works, such as [16] [17] [18] [19].

Therefore, we are currently focused on evaluating the cost of various key exchange configurations. The chosen key exchange determines various security properties of the communication, namely the presence or absence of PFS and authentication guarantee. Each one of those security services has an associated cost.

*mbedtls* comes with 161 ciphersuites [20]. Evaluating and analyzing the performance of each one manually would be very time consuming, if not impossible, given the limited time available. The fact every ciphersuite can have different configuration options (such as the certificates used and the information contained in them) aggravates this issue even more.

For this reason, we developed an automated profiling tool [21]. This tool runs a client/server connection for the provided list of ciphersuites, collects the number of CPU cycles used by the user-specified functions and presents the results in a graph. Using this tool, we were able to evaluate all of the 161 ciphersuites and reach some initial conclusions.

The tool uses *valgrind* to profile the ciphersuites and obtain the number of CPU cycles used by each one. It is important to note that *valgrind* does not count the actual number of CPU cycles used, but rather an estimate of them. All of the executables have been compiled, ran and profiled on an *x86\_64 GNU/Linux* machine (Kernel 4.15.13-1-ARCH).

The result of profiling all of the 161 *mbedtls* ciphersuites is presented in Figures 2 (for the client) and 3 (for the server). The various key exchange groups are identified by the letters in




Fig. 3. TLS Server ciphersuite profiling

red. The associated numbers refer to the approximate number of CPU cycles (in millions) used by the ciphersuite during the handshake.

For example, Figure 3 shows that the server would use 133 million CPU cycles to perform a TLS handshake using the *ECDHE-ECDSA* ciphersuite (ephemeral ECDH with ECDSA signatures).

The comparison of the *ECDHE-RSA* and *ECDH-RSA* key exchange, shows that the use of perfect forward secrecy leads to about 6.5 times more CPU cycles used during the handshake.

The results show different numbers for the client and the server. This may come down to various factors, depending on the ciphersuite. In ciphersuites that public certificates are used, during the handshake, the client needs to verify the certificate chain provided by the server, while the server does not. This means that the client will need to perform public key cryptography operations, which use a lot of CPU cycles. In the *RSA* key exchange, the server encrypts a secret and sends it to the client. The client then decrypts that secret. Since the public exponent of the *RSA* keypair is usually chosen to be small (in the case of the tested ciphersuite it's 65537), public key operations are much faster than private key operations. Since the client uses the public key to decrypt the secret, while the server encrypts it with the private key, the latter will use significantly more CPU cycles.

Ciphersuites which do not use public key cryptography (i.e. no certificates) show very similar results in both, the client and the server. Those ciphersuites are all of the pre-shared-key (PSK) ones, with the exception of *RSA-PSK*, which uses *RSA* and certificates for authentication.

As expected, PSK ciphersuites show the best results. This is why they are so widely used in embedded devices.

We performed similar measurements for the number of CPU cycles used during data encryption and decryption, as well as MAC generation and verification for each one of the ciphersuites. The scenario was a 13 MB file transfer between the client and the server. In this case, just as expected, the majority of the processing was spent on encrypting and encrypting the data. The handshake became negligible.

## VI. FUTURE WORK

Throughout April 2018 the cost of each security service will be determined more accurately. Throughout March 2018, the remaining TLS features will be profiled and a cost associated with them. This will help us understand which additional costs or savings each one of them has. Throughout March and June 2018 other relevant metrics, such as power consumption, will be collected and the the profiling will be performed on typical IoT hardware. During this same period, instead of using *valgrind* CPU cycle estimates, a transition to a tool like *PAPI* [22] will be done to obtain precise values. The automated profiling tool will be adapted accordingly. During the month of June 2018 any troubles that arise during the previous phases will be worked on.

From June until the end of August 2018, the tool that suggests a TLS configuration based on the environment's requirements and limitations will be developed. It will include a Graphical User Interface.

From mid-July until mid-September 2018 the focus will be on writing the dissertation's text.

## VII. CONCLUSION

The lack of security in IoT is a serious issue that can lead to a high monetary costs, when botnets infect the devices. Recent attacks clearly show that serious damage can be caused. An old saying attributed to the US National Security Agency (NSA) states that "Attacks always get better; they never get worse". Combined with the fact that the number of IoT devices is growing at a high pace, without any major improvements to their security, makes it clear that it is fundamental for this issue to be addressed.

While there are well established security solutions, not all of them can be used with IoT devices, due their constrained nature. One such example is the (D)TLS protocol, that because of its heavyweight nature is not suitable for a large part of IoT devices. With the proposed work, we want to contribute to this area, by designing a solution that is suitable for the IoT devices, transparent to the programmer and provides security services adaptable to the specific context needs.

## ACKNOWLEDGMENTS

I would like to thank my supervisors, Ricardo Chaves and Aleksandar Ilic for their continuous support, ideas and help that they have given me during this work.

## REFERENCES

- [1] SMALT - The Smart Home Device That Elevates Your Dining Experience. Electronic Gadget, 2017. (Accessed on 01/05/2018).
- [2] Manos Antonakakis, Tim April, Michael Bailey, Matt Bernhard, Elie Bursztein, Jaime Cochran, Zakir Durumeric, J Alex Halderman, Luca Invernizzi, Michalis Kallitsis, et al. Understanding the mirai botnet. 2017.
- [3] Brian Krebs. Reaper: Calm before the iot security storm? krebs on security. <https://krebsonsecurity.com/2017/10/reaper-calm-before-the-iot-security-storm/>, 10 2017. (Accessed on 01/04/2018).
- [4] Pull requests by iluxonchik. <https://github.com/tlswg/tls13-spec/pulls?utf8=>(Accessed on 04/12/2018).

- [5] add messages that are required to have the supported\_versions extension by iluxonchik. <https://github.com/tlswg/tls13-spec/pull/1106>. (Accessed on 04/12/2018).
- [6] fix backward compatibility section paragraph by iluxonchik. <https://github.com/tlswg/tls13-spec/pull/1122>. (Accessed on 04/12/2018).
- [7] Ssl library mbed tls / polarssl. <https://tls.mbed.org/>. (Accessed on 04/12/2018).
- [8] update test rsa certificate to use sha-256 instead of sha-1 by iluxonchik pull request #1520 armmbed/mbedtls. <https://github.com/ARMmbed/mbedtls/pull/1520>. (Accessed on 04/12/2018).
- [9] Tls-ecdh-rsa-\* ciphersuites allow ecdsa signed certificates issue #1561 armmbed/mbedtls. <https://github.com/ARMmbed/mbedtls/issues/1561>. (Accessed on 04/12/2018).
- [10] T. Dierks and E. Rescorla. The Transport Layer Security (TLS) Protocol Version 1.2. RFC 5246, RFC Editor, August 2008. <http://www.rfc-editor.org/rfc/rfc5246.txt>.
- [11] Eric Rescorla. The Transport Layer Security (TLS) Protocol Version 1.3. Internet-Draft draft-ietf-tls-tls13-26, IETF Secretariat, March 2018. <http://www.ietf.org/internet-drafts/draft-ietf-tls-tls13-26.txt>.
- [12] Carsten Bormann, Simon Lemay, Hannes Tschofenig, Klaus Hartke, Bill Silverajan, and Brian Raymor. CoAP (Constrained Application Protocol) over TCP, TLS, and WebSockets. Internet-Draft draft-ietf-core-coap-tcp-tls-11, IETF Secretariat, December 2017. <http://www.ietf.org/internet-drafts/draft-ietf-core-coap-tcp-tls-11.txt>.
- [13] C. Bormann and Z. Shelby. Block-Wise Transfers in the Constrained Application Protocol (CoAP). RFC 7959, RFC Editor, August 2016.
- [14] Shahid Raza, Ludwig Seitz, Denis Sitenkov, and Göran Selander. S3k: Scalable security with symmetric keysdtls key establishment for the internet of things. *IEEE Transactions on Automation Science and Engineering*, 13(3):1270–1280, 2016.
- [15] mbed TLS (formerly PolarSSL): SSL/TLS Library For Embedded Devices. (Accessed on 12/19/2017).
- [16] Shaify Kansal and Meenakshi Mittal. Performance evaluation of various symmetric encryption algorithms. In *Parallel, Distributed and Grid Computing (PDGC), 2014 International Conference on*, pages 105–109. IEEE, 2014.
- [17] Shaza D Riham, Ahmed Khalid, and Saife Eldin F Osman. A performance comparison of encryption algorithms aes and des. *International Journal of Engineering Research & Technology (IJERT)*, 4(12):151–154, 2015.
- [18] Priyadarshini Patil, Prashant Narayankar, DG Narayan, and SM Meena. A comprehensive evaluation of cryptographic algorithms: Des, 3des, aes, rsa and blowfish. *Procedia Computer Science*, 78:617–624, 2016.
- [19] Ram Krishna Dahal, Jagdish Bhatta, and Tanka Nath Dhamala. Performance analysis of sha-2 and sha-3 finalists. *International Journal on Cryptography and Information Security (IJCIS)*, 3(3):720–730, 2013.
- [20] Supported ssl / tls ciphersuites - mbed tls (previously polarssl). <https://tls.mbed.org/supported-ssl-ciphersuites>. (Accessed on 04/12/2018).
- [21] mbedTLS Ciphersuite Profiling Tool. <https://github.com/iluxonchik/callgrind-cpu-cycles-counter>. (Accessed on 04/12/2018).
- [22] PAPI - Performance Application Programming Interface. <http://icl.cs.utk.edu/papi/>. (Accessed on 04/14/2018).